



Stage I: Full-parameter SFT

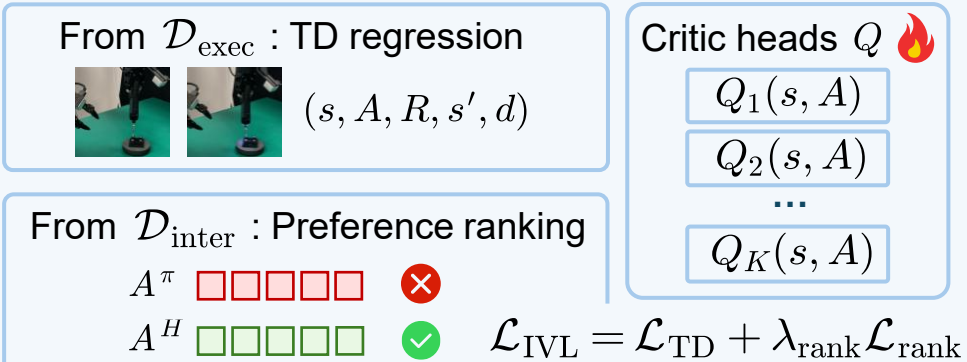


Stage II: Human-assisted online RL



Stage III: Alternating learner updates

(a) Corrective Value Learning (CVL)



(b) Adjoint-based Policy Optimization (APO)

